

Session 2: Features and Robustness Improvements

Date: 2026-01-28

Commit: d354192

Files Changed: 35 files, 6,365 lines added

Summary

Session 2 addressed high-priority gaps in Features & Robustness, implementing 8 comprehensive frameworks that significantly improve the STUNIR project's production-readiness.

1. Logging Framework (tools/logging/)

Features

- **Structured JSON logging** with RFC 8785-compatible output
- **Colored console formatter** with auto-detection of TTY
- **Log levels:** DEBUG, INFO, WARNING, ERROR, CRITICAL
- **Custom handlers:** Rotating file, syslog, async
- **Context managers** for scoped logging
- **Filters:** Level, context, sampling, rate-limiting

Usage Example

```
from tools.logging import get_logger, configure_logging, log_context

# Configure logging
configure_logging(level='INFO', json_output=True)

# Get logger
logger = get_logger('myapp')

# Log with context
with log_context(request_id='abc123', user='admin'):
    logger.info('Processing request')
```

2. Configuration Management (tools/config/)

Features

- **Multiple sources:** YAML, JSON, TOML files
- **Environment variables** with nested key support
- **CLI arguments** integration
- **Priority-based merging** of configurations

- **Schema validation** with custom validators
- **Hot-reload** support for config changes

Usage Example

```
from tools.config import ConfigManager, ConfigSchema, Field

# Create manager
manager = ConfigManager()
manager.load_file('config.yaml', priority=50)
manager.load_env(prefix='STUNIR_')

# Get merged config
config = manager.get_config()
print(config.database.host) # Dot-notation access
```

3. Retry Logic (tools/retry/)

Features

- **Retry decorator** for functions
- **Backoff strategies**: Exponential, linear, constant, jittered, Fibonacci
- **Circuit breaker** pattern for cascading failure prevention
- **Retry policies**: Max attempts, timeout, conditional
- **Async support** for coroutines
- **Statistics tracking**

Usage Example

```
from tools.retry import retry, CircuitBreaker, ExponentialBackoff

# Decorator usage
@retry(max_attempts=5, backoff=ExponentialBackoff(base=1.0))
def call_api():
    return requests.get('https://api.example.com')

# Circuit breaker
breaker = CircuitBreaker(failure_threshold=5, recovery_timeout=30)

@breaker
def call_service():
    return external_service.call()
```

4. Caching Layer (tools/common/cache_backends.py)

Features

- **Memory backend**: LRU with TTL
- **File backend**: Persistent cache with pickle
- **Redis backend**: Optional with fallback

- **Tiered caching:** L1 (memory) + L2 (file/redis)
- **Cache decorators** for functions
- **Cache warming** utilities

Usage Example

```
from tools.common.cache_backends import cached, MemoryBackend, TieredCache

# Simple caching
@cached(backend=MemoryBackend(maxsize=100), ttl=300)
def expensive_computation(x):
    return compute(x)

# Tiered cache
cache = TieredCache(
    l1=MemoryBackend(maxsize=100),
    l2=FileBackend(cache_dir='.cache')
)
```

5. Telemetry (tools/telemetry/)

Features

- **Counters:** Monotonically increasing values
- **Gauges:** Up/down values
- **Histograms:** Distribution tracking with buckets
- **Timers:** Operation duration measurement
- **Exporters:** Prometheus, JSON, StatsD
- **Collectors:** System and process metrics

Usage Example

```
from tools.telemetry import get_registry, Timer, PrometheusExporter

registry = get_registry()
counter = registry.counter('requests_total')
timer = registry.timer('request_duration')

@timer
def handle_request():
    counter.inc()
    # Process request

# Export metrics
exporter = PrometheusExporter()
print(exporter.export(registry.collect_all()))
```

6. Graceful Degradation (tools/resilience/)

Features

- **Fallback handlers** with default values
- **Fallback chains** for multiple backup options
- **Health checks**: Function, dependency, registry
- **Bulkhead pattern** for isolation
- **Graceful shutdown** with signal handling
- **Partial failure handling** for batch operations

Usage Example

```
from tools.resilience import Fallback, FallbackChain, HealthRegistry

# Fallback
@Fallback(default_value=[])
def fetch_users():
    return api.get_users()

# Health checks
registry = HealthRegistry()
registry.register(FunctionHealthCheck('database', lambda: db.ping()))
report = registry.get_report()
```

7. Rate Limiting (tools/ratelimit/)

Features

- **Token bucket**: Burst-capable rate limiting
- **Leaky bucket**: Smooth rate limiting
- **Fixed window**: Per-period limits
- **Sliding window**: Accurate per-period limits
- **Decorators**: `@rate_limit`, `@throttle`
- **Statistics tracking**

Usage Example

```
from tools.ratelimit import rate_limit, TokenBucket

# Decorator
@rate_limit(rate=10, capacity=50) # 10/sec with burst of 50
def api_call():
    return requests.get(url)

# Manual limiter
limiter = TokenBucket(rate=100, capacity=200)
if limiter.acquire():
    process_request()
```

8. Resource Management (tools/resources/)

Features

- **Resource pools:** Generic, connection, thread
- **Memory limits:** Check and enforce
- **Time limits:** Operation timeouts
- **Lifecycle management:** Register, cleanup
- **Resource monitoring:** Usage tracking
- **Lazy resources:** On-demand initialization

Usage Example

```
from tools.resources import ResourcePool, ResourceManager, enforce_limits

# Resource pool
pool = ResourcePool(factory=create_connection, max_size=10)
with pool.connection() as conn:
    conn.execute('SELECT 1')

# Resource limits
@enforce_limits(max_memory_mb=100, max_time_seconds=30)
def process_data(data):
    return heavy_computation(data)
```

File Structure

```

tools/
└── logging/
    ├── __init__.py
    ├── logger.py      # Main logging interface
    ├── formatters.py # JSON, colored formatters
    ├── handlers.py   # File rotation, syslog, async
    ├── filters.py    # Level, context, sampling
    └── context.py    # Log context management
    └── config/
        ├── __init__.py
        ├── config.py    # Configuration manager
        ├── loaders.py   # YAML, JSON, TOML, env loaders
        ├── validators.py # Schema validation
        └── schema.py    # Field definitions
    └── retry/
        ├── __init__.py
        ├── retry.py     # Retry decorator
        ├── backoff.py   # Backoff strategies
        ├── policies.py  # Retry policies
        └── circuit_breaker.py # Circuit breaker
    └── common/
        └── cache_backends.py # Multi-backend caching
    └── telemetry/
        ├── __init__.py
        ├── metrics.py    # Counter, gauge, histogram
        ├── exporters.py  # Prometheus, JSON, StatsD
        └── collectors.py # System metrics
    └── resilience/
        ├── __init__.py
        ├── fallback.py   # Fallback handlers
        ├── health.py    # Health checks
        ├── isolation.py # Bulkhead, timeout
        └── shutdown.py  # Graceful shutdown
    └── ratelimit/
        ├── __init__.py
        ├── limiter.py   # Rate limit algorithms
        └── decorators.py # @rate_limit, @throttle
    └── resources/
        ├── __init__.py
        ├── pool.py       # Resource pooling
        ├── limits.py    # Memory, time limits
        ├── lifecycle.py # Lifecycle management
        └── monitor.py   # Resource monitoring

```

Testing

All features include comprehensive tests in `tests/tools/test_session2_features.py` :

- Unit tests for each component
- Integration tests combining features
- Smoke tests for imports and basic functionality

Impact

These improvements provide STUNIR with:

- **Production-ready logging** for debugging and monitoring
- **Flexible configuration** for different environments
- **Resilient operations** with retry and circuit breaker
- **Performance optimization** through caching
- **Observability** via telemetry and metrics
- **Stability** through graceful degradation
- **Protection** via rate limiting
- **Efficiency** through resource management